

Pre-Training and Fine-Tuning Large Language Models via Unidirectional Modeling for Text Generation

Ruijun Feng

Program Analysis Group

University of New South Wales

03/07/2024

What is large language model

- Large language model (LLM) is an artificial neural network, usually with **millions or billions of parameters**, trained on **vast amounts of text data** to understand and generate human-like language.

Unidirectional modelling

- The LLM only see text in **one direction**.

Unidirectional modelling

- The LLM only see text in **one direction**.
- **Predict the next word** in a sequence based only on the **preceding context**.

Unidirectional modelling

- The LLM only see text in **one direction**.
- **Predict the next word** in a sequence based only on the **preceding context**.

"The cat sits on the mat." → "The cat [MASK]"

Unidirectional modelling

- The LLM only see text in **one direction**.
- **Predict the next word** in a sequence based only on the **preceding context**.
- Famous models: GPT, Llama, etc.

Focus of this presentation

- Basic components used in LLM, including **Tokenization**, **Embedding layer**, and **Self-Attention layer**.

Focus of this presentation

- Basic components used in LLM, including **Tokenization**, **Embedding layer**, and **Self-Attention layer**.
- Process of Pre-Training and Fine-Tuning LLM via **Teacher Forcing**.

Focus of this presentation

- Basic components used in LLM, including **Tokenization**, **Embedding layer**, and **Self-Attention layer**.
- Process of Pre-Training and Fine-Tuning LLM via **Teacher Forcing**.
- Examples based on GPT2.

Tokenization

- The purpose of tokenization is to convert **strings into integers**, so the LLM can read them.

Tokenization

- Tokenization will cut a **sentence** into **tokens**, then use a **dictionary (vocabulary)** to store the mapping relationships between **tokens** and their corresponding **integer ids**.

Tokenization

- Suppose we are using code data to pre-train an LLM, the code fragment looks like this:

```
def add(a, b):  
    return a + b
```

Tokenization

- The string format with "\\n" of the code fragment:

```
def add(a, b):\n    return a + b
```

Tokenization

- The string format with "\\n" of the code fragment:

```
def add(a, b):\n    return a + b
```

- After tokenization via GPT2 ("_" means space):

```
["def", "_add", "(", "a", ",", "_b", "):", "\\n", "_", "_",  
 "_", "_return", "_a", "_+", "_b"]
```

Tokenization

- Vocabulary for each token:

```
{"def": 4299, "_add": 751, "(": 7, "a": 64, ",": 11, "_b": 275, "):": 2599, "\n": 198, "_": 220, "_return": 1441, "_a": 257, "_+": 1343}
```

Tokenization

- The string format with escape characters "\\n" of the code fragment:

```
def add(a, b):\n    return a + b
```

- After tokenization via GPT2 ("_" means space):

```
["def", "_add", "(", "a", ",", "b", "):", "\\n", "_", "_",  
 "_", "_return", "_a", "_+", "_b"]
```

- Integer ids:

```
[4299, 751, 7, 64, 11, 275, 2599, 198, 220, 220, 220,  
1441, 257, 1343, 275]
```


Tokenization

- To encode more semantic relationships, each integer id will be converted into an **embedding vector via embedding layer.**

Embedding layer

- What is embedding layer in LLM?

Embedding layer

- Embedding layer is a **list of trainable vectors**. Each vectors have the **same dimensions**.

vector 0: $[0.1, -0.2, 0.3, 0.4, \dots, 0.6]$,

vector 1: $[0.7, -0.8, 0.9, -1.0, \dots, -1.2]$,

vector 2: $[-1.3, 1.4, -1.5, 1.6, \dots, 1.8]$,

$\dots \dots,$

vector n : $[2.5, -2.6, 2.7, -2.8, \dots, -3.0]$

Embedding layer


- There are two types of embedding layers used in GPT2:
 - Word embedding layer
 - Positional embedding layer

Word embedding layer

- Integer id is the index, use it to index the list in the word embedding layer.
- The corresponding vector is the vectorized representation of that token.

Token: #

Integer id: 2



vector 0: [0.1, -0.2, 0.3, 0.4, ..., 0.6],

vector 1: [0.7, -0.8, 0.9, -1.0, ..., -1.2],

vector 2: [-1.3, 1.4, -1.5, 1.6, ..., 1.8],

... ..,

vector n : [2.5, -2.6, 2.7, -2.8, ..., -3.0]]

Word embedding layer

- "**def**" is a token, its integer id is "**4299**".
- The 4299-th vector of the word embedding layer is the word embedding vector of "**def**".

```
• # Get the embedding vector of "def"  
• word_embedding =  
  word_embedding_layer(torch.tensor([4299]))
```

Word embedding layer

```
# Declare a word embedding layer
word_embedding_layer = torch.nn.Embedding(50257, 768)
# Initialize a token ids tensor
token_ids = torch.tensor([4299, 751, 7, 64, 11, 275,
2599, 198, 220, 220, 220, 1441, 257, 1343, 275])
# Get the embedding vectors of the token ids
word_embeddings = word_embedding_layer(token_ids)
```

Word embedding layer

	Token	Token ID	Word Embedding Vector
0	def	4299	<code>[-0.075, -0.083, 0.146, ..., -0.057]</code>
1	_add	751	<code>[0.129, -0.006, 0.129, ..., 0.111]</code>
2	(7	<code>[-0.130, -0.212, 0.132, ..., -0.075]</code>
3	a	64	<code>[-0.176, -0.099, 0.206, ..., -0.055]</code>
4	,	11	<code>[0.011, -0.003, 0.032, ..., -0.060]</code>
5	_b	275	<code>[0.047, -0.018, 0.062, ..., 0.132]</code>
6):	2599	<code>[-0.080, -0.217, -0.003, ..., 0.116]</code>
7	\n	198	<code>[-0.001, 0.018, 0.053, ..., -0.035]</code>
8	_	220	<code>[0.096, -0.091, 0.085, ..., 0.126]</code>
9	_	220	<code>[0.096, -0.091, 0.085, ..., 0.126]</code>
10	_	220	<code>[0.096, -0.091, 0.085, ..., 0.126]</code>
11	_return	1441	<code>[0.029, -0.021, 0.103, ..., -0.124]</code>
12	_a	257	<code>[-0.051, 0.006, 0.047, ..., -0.038]</code>
13	_+	1343	<code>[-0.045, -0.093, 0.013, ..., -0.112]</code>
14	_b	275	<code>[0.047, -0.018, 0.062, ..., 0.132]</code>

Word embedding layer

- Word embedding layer **cannot describe the relative positions** between different tokens.

Word embedding layer

	Token	Token ID	Word Embedding Vector
0	def	4299	[-0.075, -0.083, 0.146, ..., -0.057]
1	_add	751	[0.129, -0.006, 0.129, ..., 0.111]
2	(7	[-0.130, -0.212, 0.132, ..., -0.075]
3	a	64	[-0.176, -0.099, 0.206, ..., -0.055]
4	,	11	[0.011, -0.003, 0.032, ..., -0.060]
5	_b	275	[0.047, -0.018, 0.062, ..., 0.132]
6):	2599	[-0.080, -0.217, -0.003, ..., 0.116]
7	\n	198	[-0.001, 0.018, 0.053, ..., -0.035]
8	_	220	[0.096, -0.091, 0.085, ..., 0.126]
9	_	220	[0.096, -0.091, 0.085, ..., 0.126]
10	_	220	[0.096, -0.091, 0.085, ..., 0.126]
11	_return	1441	[0.029, -0.021, 0.103, ..., -0.124]
12	_a	257	[-0.051, 0.006, 0.047, ..., -0.038]
13	+	1343	[-0.045, -0.093, 0.013, ..., -0.112]
14	_b	275	[0.047, -0.018, 0.062, ..., 0.132]

Positional embedding layer

- Positional embedding layer aims to describe the **relative positions** between different tokens using the **position index**.

Positional embedding layer

```
# Declare a position embedding layer  
position_embedding_layer = torch.nn.Embedding(1024,  
768)
```

Positional embedding layer

- List of tokens:

```
["def", "_add", "(", "a", ",", "_b", "):", "\n",  
" ", " ", " ", "_return", "_a", "_+", "_b"]
```

Positional embedding layer

- List of tokens:

```
["def", "_add", "(", "a", ",", "_b", "):", "\n",  
" ", " ", " ", "_return", "_a", "_+", "_b"]
```

- The tokenized code fragment has 15 tokens, so the position index is [0, 1, ..., 14].

Positional embedding layer

- List of tokens:

```
["def", "_add", "(", "a", ",", "_b", "):", "\n",  
" ", " ", " ", "_return", "_a", "_+", "_b"]
```

- The tokenized code fragment has 15 tokens, so the position index is [0, 1, ..., 14].
- "def" and 0, "_add" and 1, ..., "_b" and 14

Positional embedding layer

```
# Declare a position embedding layer
position_embedding_layer = torch.nn.Embedding(1024,
768)
# Initialize a position index from 0 to 14
position_index = torch.arange(15)
```


Positional embedding layer

```
# Declare a position embedding layer
position_embedding_layer = torch.nn.Embedding(1024,
768)
# Initialize a position index from 0 to 14
position_index = torch.arange(15)
# Get the embedding vectors of the position index
position_embeddings =
position_embedding_layer(position_index)
```

Positional embedding layer

	Position Index	Token	Position Embedding Vector
0	0	def	[-0.019, -0.197, 0.004, ..., 0.054]
1	1	_add	[0.024, -0.054, -0.095, ..., -0.000]
2	2	([0.004, -0.085, 0.055, ..., -0.021]
3	3	a	[-0.000, -0.074, 0.106, ..., -0.007]
4	4	,	[0.008, -0.025, 0.127, ..., -0.007]
5	5	_b	[0.010, -0.034, 0.131, ..., -0.007]
6	6):	[0.003, -0.021, 0.120, ..., -0.003]
7	7	\n	[0.003, -0.003, 0.117, ..., -0.007]
8	8	_	[-0.001, -0.002, 0.111, ..., -0.010]
9	9	_	[0.005, 0.002, 0.118, ..., -0.006]
10	10	_	[0.002, 0.006, 0.100, ..., -0.006]
11	11	_return	[-0.004, 0.017, 0.107, ..., -0.006]
12	12	_a	[0.000, 0.017, 0.097, ..., -0.008]
13	13	_+	[0.004, 0.020, 0.105, ..., 0.000]
14	14	_b	[0.001, 0.023, 0.096, ..., -0.002]

Final input embedding

- By **adding** the **word embedding vectors** and **positional embedding vectors**, we have the input embedding that encodes all token information.

```
# Combine the word and position embeddings  
final_embeddings = word_embeddings + position_embeddings
```

Final input embedding

- In GPT2, the size of vocabulary is set to **50257**.
- The maximum number of position indexes is set to **1024**.
- The dimension of the encoded embedding vector is **768**.

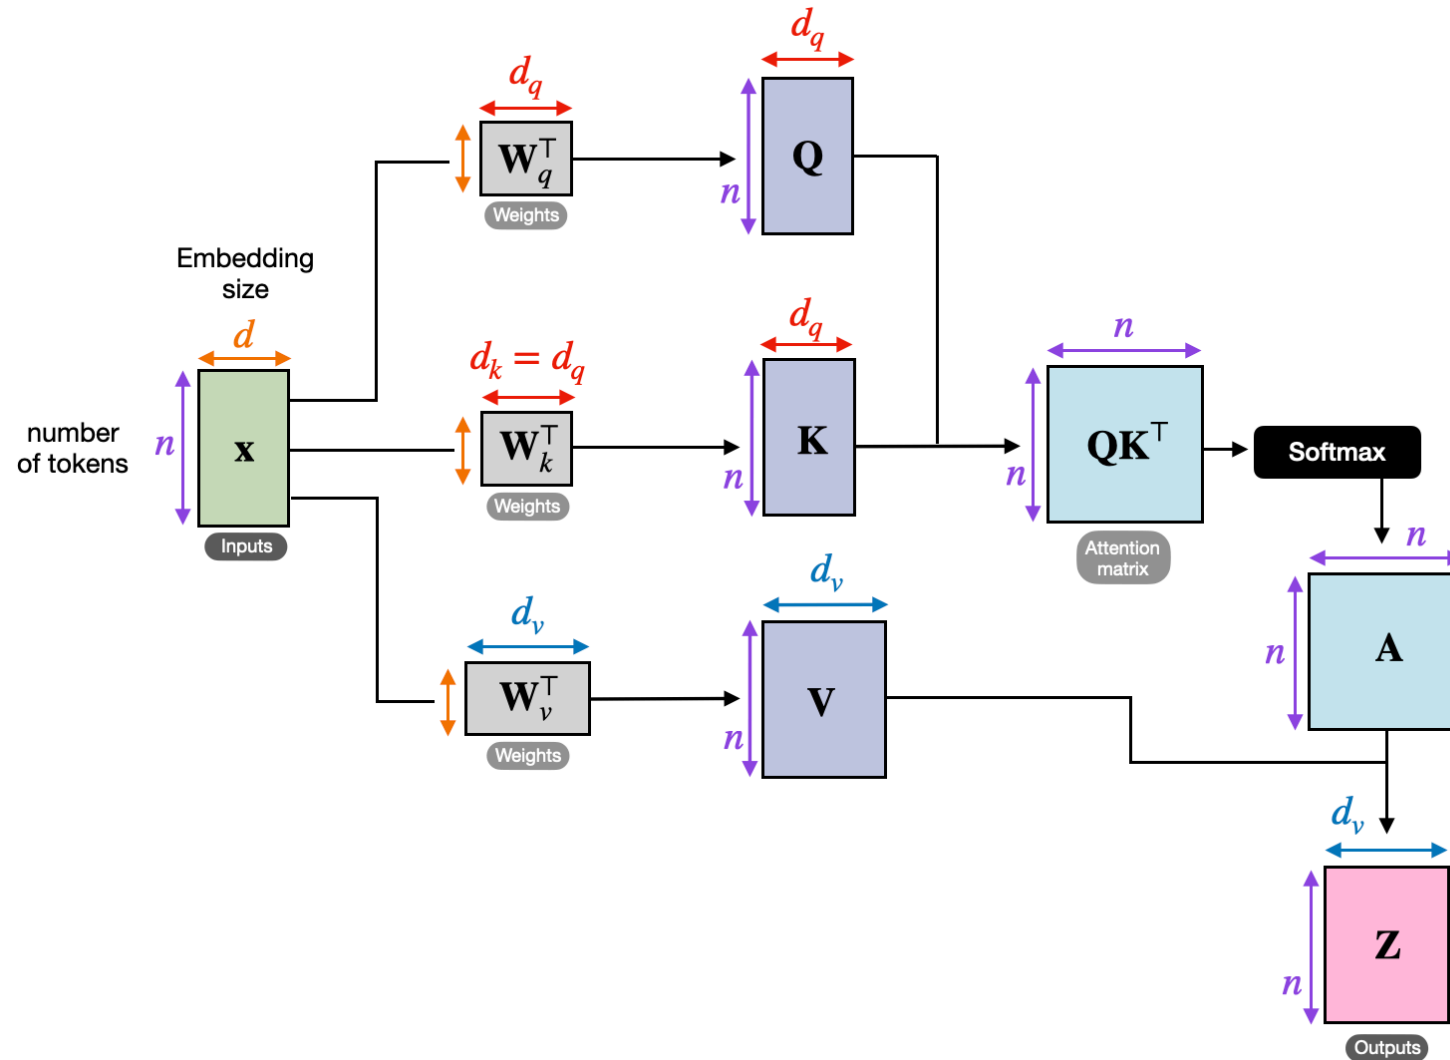
Self-attention layer

- Self-attention layer is used to calculate the relationships between different tokens using the input embedding.

Self-attention layer

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Self-attention layer

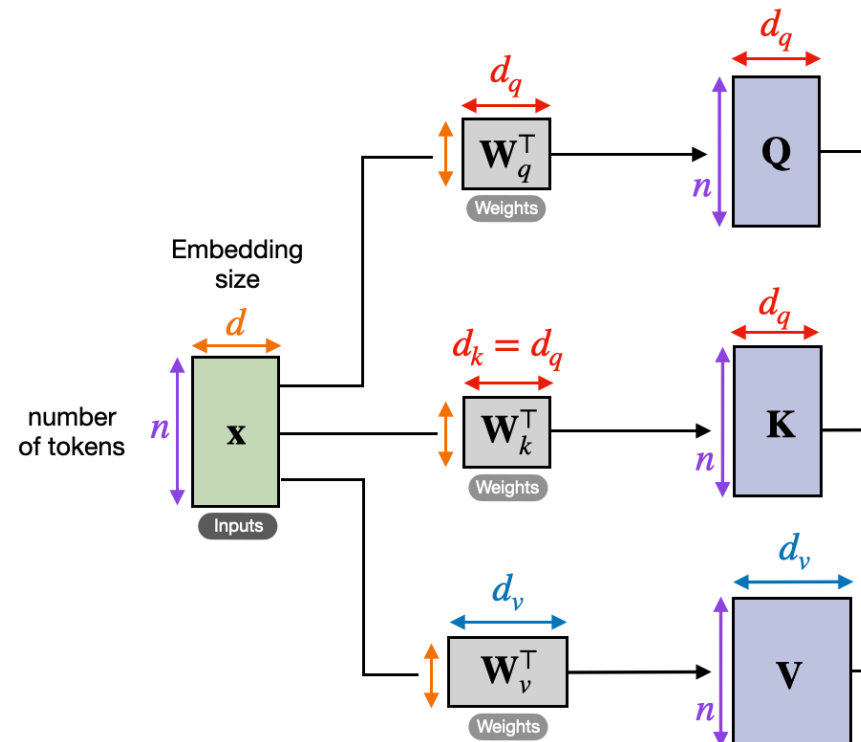


Self-attention layer—Q, K, V

$$\text{Attention}(\boxed{Q, K, V}) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Self-attention layer—Q, K, V

- First, three **separate** linear layers are used to encode the input embeddings into **hidden vectors**, denoted as **query** (Q), **key** (K), and **value** (V).



Self-attention layer—Q, K, V

- **Query (Q)** is the content that is being looked for.
- **Key (K)** is reference of the content.
- **Value (V)** is the content that is being searched.

Self-attention layer—Q, K, V

```
# Define linear layers for Q, K, V
hidden_dim = 768 # Same as the dimension of embedding vectors
linear_q = nn.Linear(hidden_dim, hidden_dim)
linear_k = nn.Linear(hidden_dim, hidden_dim)
linear_v = nn.Linear(hidden_dim, hidden_dim)
# Apply linear transformations to get Q, K, V
Q = linear_q(final_embeddings)
K = linear_k(final_embeddings)
V = linear_v(final_embeddings)
```

Self-attention layer—Q, K, V

- Each Q , K , V contains 15 vectors (corresponding to the number of tokens in the code fragment).

Token	Q (15, 768)	Token	K (15, 768)	Token	V (15, 768)
0 def	[1.188, -4.549, ..., 2.721]	0 def	[-1.083, 1.734, ..., -0.694]	0 def	[0.092, -0.522, ..., 0.589]
1 _add	[2.478, -1.816, ..., -1.107]	1 _add	[0.081, 1.011, ..., -1.701]	1 _add	[0.505, 0.107, ..., -0.431]
2 ([1.280, -1.915, ..., 0.712]	2 ([-1.008, -1.505, ..., -1.077]	2 ([-0.038, -0.236, ..., -0.596]
3 a	[1.466, 0.061, ..., -0.108]	3 a	[-1.647, -1.664, ..., -1.361]	3 a	[0.255, -0.025, ..., -0.514]
4 ,	[0.958, -0.852, ..., -0.022]	4 ,	[0.204, -0.925, ..., -1.277]	4 ,	[0.157, 0.019, ..., -0.276]
5 _b	[1.282, -1.552, ..., -0.977]	5 _b	[0.094, -1.595, ..., -1.411]	5 _b	[-0.023, -0.046, ..., -0.314]
6):	[0.879, -2.032, ..., 0.415]	6):	[0.068, -0.975, ..., -1.570]	6):	[0.220, 0.038, ..., -0.184]
7 \n	[0.315, -0.799, ..., 0.637]	7 \n	[0.296, -1.048, ..., -0.775]	7 \n	[0.169, 0.043, ..., -0.396]
8 _	[0.947, -0.437, ..., -0.532]	8 _	[-1.036, -1.265, ..., -0.690]	8 _	[0.222, 0.114, ..., -0.077]
9 _	[0.771, -0.351, ..., -0.607]	9 _	[-1.010, -1.216, ..., -0.721]	9 _	[0.231, 0.129, ..., -0.094]
10 _	[0.768, -0.359, ..., -0.653]	10 _	[-0.938, -1.202, ..., -0.740]	10 _	[0.233, 0.137, ..., -0.086]
11 _return	[0.007, -0.775, ..., 0.572]	11 _return	[0.065, -0.034, ..., -1.387]	11 _return	[0.143, 0.072, ..., -0.076]
14 _b	[0.718, -1.453, ..., -1.470]	14 _b	[0.570, -1.329, ..., -1.370]	14 _b	[-0.014, 0.054, ..., -0.275]

Self-attention layer—Q, K, V

- Each vector is 768-dimensional (corresponding to the dimension of embeddings).


Q (15, 768)			K (15, 768)			V (15, 768)		
Token			Token			Token		
0	def	[1.188, -4.549, ..., 2.721]	0	def	[-1.083, 1.734, ..., -0.694]	0	def	[0.092, -0.522, ..., 0.589]
1	_add	[2.478, -1.816, ..., -1.107]	1	_add	[0.081, 1.011, ..., -1.701]	1	_add	[0.505, 0.107, ..., -0.431]
2	([1.280, -1.915, ..., 0.712]	2	([-1.008, -1.505, ..., -1.077]	2	([-0.038, -0.236, ..., -0.596]
3	a	[1.466, 0.061, ..., -0.108]	3	a	[-1.647, -1.664, ..., -1.361]	3	a	[0.255, -0.025, ..., -0.514]
4	,	[0.958, -0.852, ..., -0.022]	4	,	[0.204, -0.925, ..., -1.277]	4	,	[0.157, 0.019, ..., -0.276]
5	_b	[1.282, -1.552, ..., -0.977]	5	_b	[0.094, -1.595, ..., -1.411]	5	_b	[-0.023, -0.046, ..., -0.314]
6):	[0.879, -2.032, ..., 0.415]	6):	[0.068, -0.975, ..., -1.570]	6):	[0.220, 0.038, ..., -0.184]
7	\n	[0.315, -0.799, ..., 0.637]	7	\n	[0.296, -1.048, ..., -0.775]	7	\n	[0.169, 0.043, ..., -0.396]
8	_	[0.947, -0.437, ..., -0.532]	8	_	[-1.036, -1.265, ..., -0.690]	8	_	[0.222, 0.114, ..., -0.077]
9	_	[0.771, -0.351, ..., -0.607]	9	_	[-1.010, -1.216, ..., -0.721]	9	_	[0.231, 0.129, ..., -0.094]
10	_	[0.768, -0.359, ..., -0.653]	10	_	[-0.938, -1.202, ..., -0.740]	10	_	[0.233, 0.137, ..., -0.086]
11	_return	[0.007, -0.775, ..., 0.572]	11	_return	[0.065, -0.034, ..., -1.387]	11	_return	[0.143, 0.072, ..., -0.076]
14	_b	[0.718, -1.453, ..., -1.470]	14	_b	[0.570, -1.329, ..., -1.370]	14	_b	[-0.014, 0.054, ..., -0.275]

Self-attention layer — Attention Scores

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Self-attention layer—Attention Scores

- Use Q and K to get attention scores matrix, which indicates the relevance of each token in the sequence (key) to the current token (query).



Self-attention layer—Attention Scores

- Calculate QK^T and get a $n \times n$ scaled attention scores matrix, where the n is the number of tokens in the code fragment (i.e., 15).

```
# Calculate the matrix multiplication between Q and K^T  
attention_scores = torch.matmul(Q, K.transpose(-2, -1))
```


Self-attention layer—Attention Scores

- First line describes the relationships of "**def**" token to other tokens.

	def	_add	(a	,	_b):	\n	_	_	_	_return	_a	_+	_b
def	1221.9	410.3	394.8	432.0	497.4	531.8	423.4	473.5	589.4	588.0	589.4	445.1	529.3	547.6	534.9
_add	339.1	555.8	471.9	424.8	378.1	420.2	374.6	328.0	294.6	285.0	279.3	350.6	277.5	297.0	341.3
(262.5	382.1	401.5	373.3	341.1	374.7	339.9	306.7	284.1	275.1	271.1	302.5	254.7	283.6	311.0
a	291.6	355.9	348.6	476.9	341.1	357.1	303.1	344.4	281.3	274.0	268.4	289.6	297.7	279.5	288.8
,	246.4	327.8	344.1	330.8	374.4	354.0	253.3	305.9	259.9	249.2	244.0	223.3	265.0	240.4	280.3
_b	338.1	345.3	351.6	377.2	385.3	442.3	324.2	342.4	323.6	313.2	307.2	287.1	313.7	276.6	369.8
):	450.5	343.7	303.7	329.4	296.9	335.4	448.7	307.3	300.0	292.2	288.3	242.0	267.1	298.9	279.2
\n	308.8	309.5	313.9	316.5	328.2	329.2	264.2	350.8	263.8	257.2	252.1	229.2	251.2	239.4	270.4
_	419.0	275.6	285.2	302.4	309.8	358.1	295.8	270.5	349.8	343.0	339.2	258.8	257.4	289.3	308.4
_	422.1	268.2	275.0	293.6	299.0	348.6	287.2	261.7	343.5	337.3	333.5	252.7	251.5	282.7	301.8
_return	456.9	301.8	277.2	314.2	284.0	317.6	267.9	309.5	316.1	310.0	306.4	342.9	252.7	264.0	271.7
_a	433.5	283.1	251.3	278.8	287.4	300.1	203.5	267.9	253.9	250.0	247.4	207.1	284.2	231.1	267.8
_+	372.1	245.7	242.7	265.9	251.2	282.2	241.0	258.1	270.4	264.5	260.4	194.8	220.5	306.2	236.4
_b	373.8	282.7	273.8	305.1	314.0	379.0	261.3	286.4	282.4	275.6	270.5	239.9	279.4	230.2	326.7

Self-attention layer—Attention Scores

- The variance of the attention scores is too big (**8936.96**).
- **Vanishing gradient problem**: the gradients become too small to effectively update the parameters during training.

```
>>> # Calculate the variance of the attention scores
>>> variance = torch.var(attention_scores)
>>> print("\nVariance of Attention Scores:", variance.item())
>>> Variance of Attention Scores: 8936.9609375
```

Self-attention layer—Scaled Attention Scores

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Self-attention layer—Scaled Attention Scores

- Calculate QK^T / \sqrt{d} to scale the attention scores matrix, where d is the hidden dim (i.e., 768) for scaling.

```
# Scale the attention scores by the square root of the hidden  
dimension  
scaled_attention_scores = attention_scores / (Q.size(-1) ** 0.5)
```

Self-attention layer—Scaled Attention Scores

	def	_add	(a	,	_b):	\n	_	_	_	_return	_a	_+	_b
def	44.1	14.8	14.2	15.6	17.9	19.2	15.3	17.1	21.3	21.2	21.3	16.1	19.1	19.8	19.3
_add	12.2	20.1	17.0	15.3	13.6	15.2	13.5	11.8	10.6	10.3	10.1	12.7	10.0	10.7	12.3
(9.5	13.8	14.5	13.5	12.3	13.5	12.3	11.1	10.3	9.9	9.8	10.9	9.2	10.2	11.2
a	10.5	12.8	12.6	17.2	12.3	12.9	10.9	12.4	10.2	9.9	9.7	10.4	10.7	10.1	10.4
,	8.9	11.8	12.4	11.9	13.5	12.8	9.1	11.0	9.4	9.0	8.8	8.1	9.6	8.7	10.1
_b	12.2	12.5	12.7	13.6	13.9	16.0	11.7	12.4	11.7	11.3	11.1	10.4	11.3	10.0	13.3
):	16.3	12.4	11.0	11.9	10.7	12.1	16.2	11.1	10.8	10.5	10.4	8.7	9.6	10.8	10.1
\n	11.1	11.2	11.3	11.4	11.8	11.9	9.5	12.7	9.5	9.3	9.1	8.3	9.1	8.6	9.8
_	15.1	9.9	10.3	10.9	11.2	12.9	10.7	9.8	12.6	12.4	12.2	9.3	9.3	10.4	11.1
_	15.2	9.7	9.9	10.6	10.8	12.6	10.4	9.4	12.4	12.2	12.0	9.1	9.1	10.2	10.9
_return	16.5	10.9	10.0	11.3	10.2	11.5	9.7	11.2	11.4	11.2	11.1	12.4	9.1	9.5	9.8
_a	15.6	10.2	9.1	10.1	10.4	10.8	7.3	9.7	9.2	9.0	8.9	7.5	10.3	8.3	9.7
_+	13.4	8.9	8.8	9.6	9.1	10.2	8.7	9.3	9.8	9.5	9.4	7.0	8.0	11.0	8.5
_b	13.5	10.2	9.9	11.0	11.3	13.7	9.4	10.3	10.2	9.9	9.8	8.7	10.1	8.3	11.8

```

>>> # Calculate the variance of the scaled attention scores
>>> variance = torch.var(scaled_attention_scores)
>>> print("\nVariance of Scaled Attention Scores:", variance.item())
>>> Variance of Scaled Attention Scores: 11.63666820526123

```

Self-attention layer—Softmax and Masking

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Self-attention layer—Softmax and Masking

- Before softmax, masking operation is used to cast the **upper triangle** as $-\infty$
- This ensure the LLM can **only see tokens from left to right**.
- Because in the text generation task, only the **preceding context** is available.

Self-attention layer—Softmax and Masking

```
# Create a mask to mask the upper triangle of the attention scores matrix
mask = torch.triu(torch.ones_like(attention_scores), diagonal=1).bool()
attention_scores_masked = attention_scores.masked_fill(mask, float('-inf'))
```

Masked Attention Scores Matrix (K x Q):

def	44.1	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf
_add	12.2	20.1	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf
a	10.5	12.8	12.6	17.2	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf
,	8.9	11.8	12.4	11.9	12.5	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf
_b	12.2	12.5	12.7	13.6	13.9	10.0	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf
):	16.3	12.4	11.0	11.9	10.7	12.1	10.2	-inf	-inf	-inf	-inf	-inf	-inf	-inf	-inf
\n	11.1	11.2	11.3	11.4	11.8	11.9	9.5	12.7	-inf	-inf	-inf	-inf	-inf	-inf	-inf
_	15.1	9.9	10.3	10.9	11.2	12.9	10.7	9.8	12.5	-inf	-inf	-inf	-inf	-inf	-inf
_	15.2	9.7	9.9	10.6	10.8	12.6	10.4	9.4	12.4	12.2	-inf	-inf	-inf	-inf	-inf
_	15.4	9.5	9.8	10.4	10.7	12.5	10.3	9.4	12.3	12.1	12.0	-inf	-inf	-inf	-inf
_return	16.5	10.9	10.0	11.3	10.2	11.5	9.7	11.2	11.4	11.2	11.1	12.1	-inf	-inf	-inf
_a	15.6	10.2	9.1	10.1	10.4	10.8	7.3	9.7	9.2	9.0	8.9	7.5	10.5	-inf	-inf
_+	13.4	8.9	8.8	9.6	9.1	10.2	8.7	9.3	9.8	9.5	9.4	7.0	8.0	11.0	-inf
_b	13.5	10.2	9.9	11.0	11.3	13.7	9.4	10.3	10.2	9.9	9.8	8.7	10.1	8.3	11.8

Self-attention layer—Softmax and Masking

Softmax equation:

$$\sigma(z_i) = \frac{e^{z_i}}{\sum_{j=1}^n e^{z_j}}$$

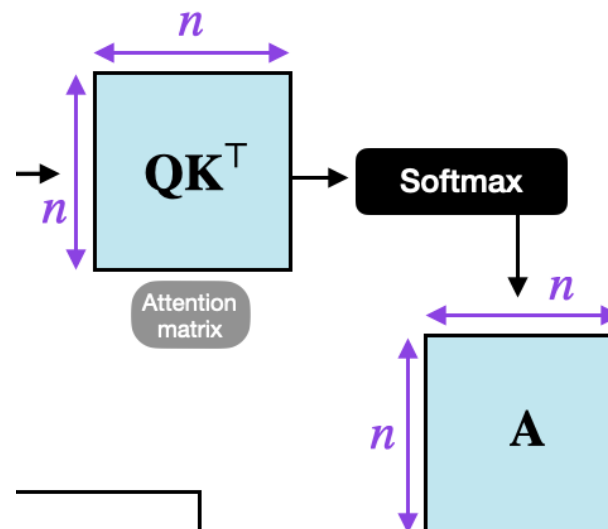
Softmax equation in masking:

$$\sigma(-\infty) = \frac{e^{-\infty}}{\sum_{j=1}^n e^{z_j}} = \frac{0}{\sum_{j=1}^n e^{z_j}} = 0$$

Self-attention layer—Attention Matrix

- Applying the softmax function to QK^T / \sqrt{d} and get the attention matrix A .

```
# Apply the softmax function to get the attention weights matrix A  
A = torch.nn.functional.softmax(attention_scores_masked, dim=-1)
```



Self-attention layer—Attention Matrix

Attention Weights Matrix A (after applying softmax):

	def	_add	(a	,	_b)	\n	_	_	_	_return	_a	_+	_b
def	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_add	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(0.0	0.3	0.7	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
a	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
,	0.0	0.1	0.2	0.1	0.6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_b	0.0	0.0	0.0	0.1	0.1	0.8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
)	0.5	0.0	0.0	0.0	0.0	0.0	0.5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
\n	0.1	0.1	0.1	0.1	0.2	0.2	0.0	0.3	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_	0.8	0.0	0.0	0.0	0.0	0.1	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0
_	0.8	0.0	0.0	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_	0.8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_return	0.9	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_a	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_+	0.8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1	0.0
_b	0.4	0.0	0.0	0.0	0.0	0.4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1

Self-attention layer—Attention Matrix

Attention Weights Matrix A (after applying softmax):

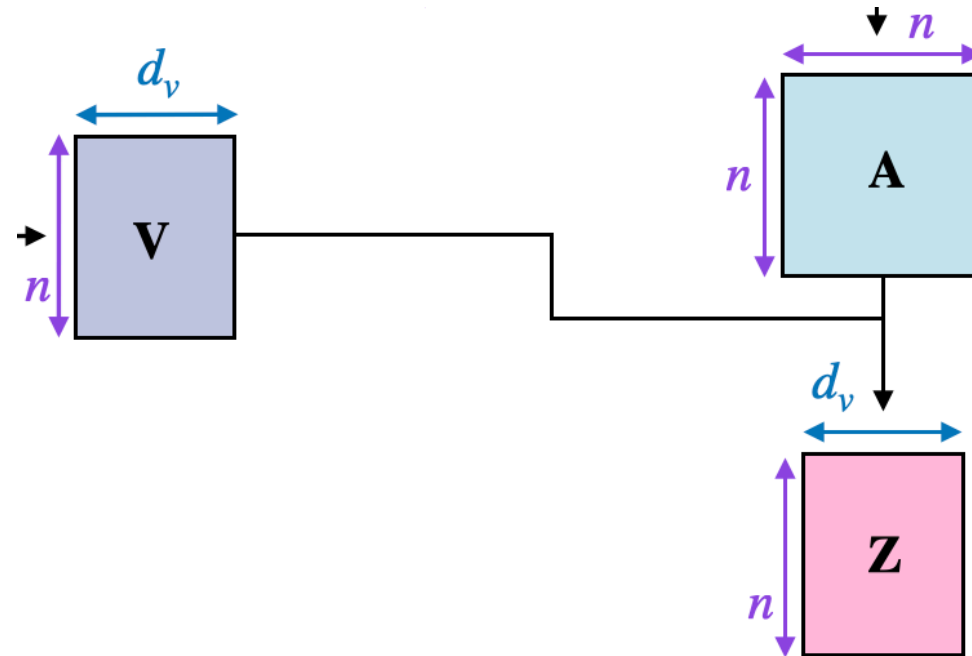
	def	_add	(a	,	_b)	\n	_	_	_	_return	_a	_+	_b
def	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_add	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(0.0	0.3	0.7	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
a	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
,	0.0	0.1	0.2	0.1	0.6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_b	0.0	0.0	0.0	0.1	0.1	0.8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
)	0.5	0.0	0.0	0.0	0.0	0.0	0.5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
\n	0.1	0.1	0.1	0.1	0.2	0.2	0.0	0.3	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_	0.8	0.0	0.0	0.0	0.0	0.1	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0
_	0.8	0.0	0.0	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_	0.8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_return	0.9	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_a	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
_+	0.8	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1	0.0
_b	0.4	0.0	0.0	0.0	0.0	0.4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1

Self-attention layer— $Z=AV$

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Self-attention layer— $Z=AV$

- Perform matrix multiplication between attention matrix A and the value matrix V and get the output $Z=AV$.
- Z has a shape of $n \times d$ (number of tokens \times hidden dim), which is 15×768 .



Self-attention layer— $Z=AV$

```
# Compute the matrix multiplication between A and V
Z = torch.matmul(A, V)
```

Each line of Z is a new presentation of the token.

	Token	Z (15, 768)
0	def	[0.092, -0.522, ..., 0.589]
1	_add	[0.505, 0.107, ..., -0.431]
2	([0.142, -0.124, ..., -0.536]
3	a	[0.255, -0.026, ..., -0.512]
4	,	[0.168, -0.029, ..., -0.378]
5	_b	[0.028, -0.048, ..., -0.319]
6):	[0.156, -0.243, ..., 0.196]
7	\n	[0.148, -0.044, ..., -0.320]
8	_	[0.096, -0.412, ..., 0.416]
9	_	[0.101, -0.421, ..., 0.446]
10	_	[0.103, -0.424, ..., 0.459]
12	_a	[0.095, -0.504, ..., 0.558]
13	_+	[0.103, -0.406, ..., 0.387]
14	_b	[0.056, -0.195, ..., 0.004]

Prediction—Output Logits

The size of vector is not the same as vocabulary.

Cannot predict the next token from vocabulary.

	Token	Z (15, 768)
0	def	[0.092, -0.522, ..., 0.589]
1	_add	[0.505, 0.107, ..., -0.431]
2	([0.142, -0.124, ..., -0.536]
3	a	[0.255, -0.026, ..., -0.512]
4	,	[0.168, -0.029, ..., -0.378]
5	_b	[0.028, -0.048, ..., -0.319]
6):	[0.156, -0.243, ..., 0.196]
7	\n	[0.148, -0.044, ..., -0.320]
8	_	[0.096, -0.412, ..., 0.416]
9	_	[0.101, -0.421, ..., 0.446]
10	_	[0.103, -0.424, ..., 0.459]
12	_a	[0.095, -0.504, ..., 0.558]
13	_+	[0.103, -0.406, ..., 0.387]
14	_b	[0.056, -0.195, ..., 0.004]

Prediction—Output Logits

- Perform another linear transformation on Z to map each vectors from hidden dim (i.e., 768) into the size of vocabulary (i.e., 50257), which is the **output logits**.
- Output logit indicates which integer id in the vocabulary the predicted token belongs to.

Self-attention layer—Output Logits

```
# Declare the linear layer for the output projection to  
the vocabulary size  
linear_output = nn.Linear(hidden_dim, 50257)  
output_logits = linear_output(Z)
```

Self-attention layer—Output Logits

- Each element of this 50257-dimensional vector indicates the probability that the predicted next token should be.

```
Output Logit (15, 50257)
0 [-3.332, 0.474, -5.038, ..., -3.483]
1 [0.346, 0.275, 0.025, ..., 0.331]
2 [0.625, 0.793, -0.324, ..., -0.682]
3 [0.919, 2.356, 1.432, ..., 2.054]
4 [1.450, 1.957, 0.323, ..., 0.154]
5 [1.076, 1.328, 0.089, ..., -0.479]
6 [-1.479, 0.659, -2.643, ..., -1.749]
7 [0.579, 1.465, 0.031, ..., -0.233]
8 [-2.470, 0.648, -3.941, ..., -2.760]
9 [-2.607, 0.620, -4.076, ..., -2.822]
10 [-2.665, 0.608, -4.131, ..., -2.845]
11 [-3.118, 0.513, -4.745, ..., -3.286]
12 [-3.191, 0.507, -4.854, ..., -3.365]
13 [-2.445, 0.601, -3.867, ..., -2.656]
14 [-0.522, 0.980, -1.689, ..., -1.553]
```

Teacher forcing

- Teacher forcing trains the LLM in a self-supervised manner by using the output logits to **predict the next token**.

Teacher forcing

- For example, the first input token is "**def**".

	Input Token	Expected Next Token	Expected Next Token ID	Output Logit (15, 50257)
0	def	_add	751	[-3.332, 0.474, -5.038, ..., -3.483]
1	_add	(7	[0.346, 0.275, 0.025, ..., 0.331]
2	(a	64	[0.625, 0.793, -0.324, ..., -0.682]
3	a	,	11	[0.919, 2.356, 1.432, ..., 2.054]
4	,	_b	275	[1.450, 1.957, 0.323, ..., 0.154]
5	_b):	2599	[1.076, 1.328, 0.089, ..., -0.479]
6):	\n	198	[-1.479, 0.659, -2.643, ..., -1.749]
7	\n	_	220	[0.579, 1.465, 0.031, ..., -0.233]
8	_	_	220	[-2.470, 0.648, -3.941, ..., -2.760]
9	_	_	220	[-2.607, 0.620, -4.076, ..., -2.822]
10	_	_return	1441	[-2.665, 0.608, -4.131, ..., -2.845]

Teacher forcing

- For example, the first input token is "**def**". LLM outputs the next token probability distribution of "**def**" as **output logit**.

	Input Token	Expected Next Token	Expected Next Token ID	Output Logit (15, 50257)
0	def	_add	751	[-3.332, 0.474, -5.038, ..., -3.483]
1	_add	(7	[0.346, 0.275, 0.025, ..., 0.331]
2	(a	64	[0.625, 0.793, -0.324, ..., -0.682]
3	a	,	11	[0.919, 2.356, 1.432, ..., 2.054]
4	,	_b	275	[1.450, 1.957, 0.323, ..., 0.154]
5	_b):	2599	[1.076, 1.328, 0.089, ..., -0.479]

Teacher forcing

- For example, the first input token is "**def**". LLM outputs the next token probability distribution of "**def**" as **output logit**. The highest value in the output logit is the predicted next token. It should be the token after "**def**", which is "**_add**" and its id is 751.

	Input Token	Expected Next Token	Expected Next Token ID	Output Logit (15, 50257)
0	def	_add	751	[-3.332, 0.474, -5.038, ..., -3.483]
1	_add	(7	[0.346, 0.275, 0.025, ..., 0.331]

Teacher forcing

- Calculate the cross-entropy loss between the output logits and the expected next token ids to update the embedding layers and the self-attention layers.

```
# Calculate the cross-entropy loss
loss =
torch.nn.functional.cross_entropy(output_logits[0:14, :].view(-1, 50257), token_ids[1:15].view(-1))
loss.backward()
```


Teacher forcing

- For the last row, it predicts the token after "**_b**", which does not have any valid next token. So, it does not participate in loss calculation.

No valid next token

`def add(a, b):\n return a + b`



	Input	Token	Expected	Next	Token	Expected	Next	Token	ID	Output	Logit (15, 50257)
14		_b			None			None			[-0.522, 0.980, -1.689, ..., -1.553]

Pre-training


- After pre-training, the LLM have learned lots of knowledge and capable of generating text based on the given context.

Supervised fine-tuning

- However, the answer provided by LLM may not satisfy the need of the user.

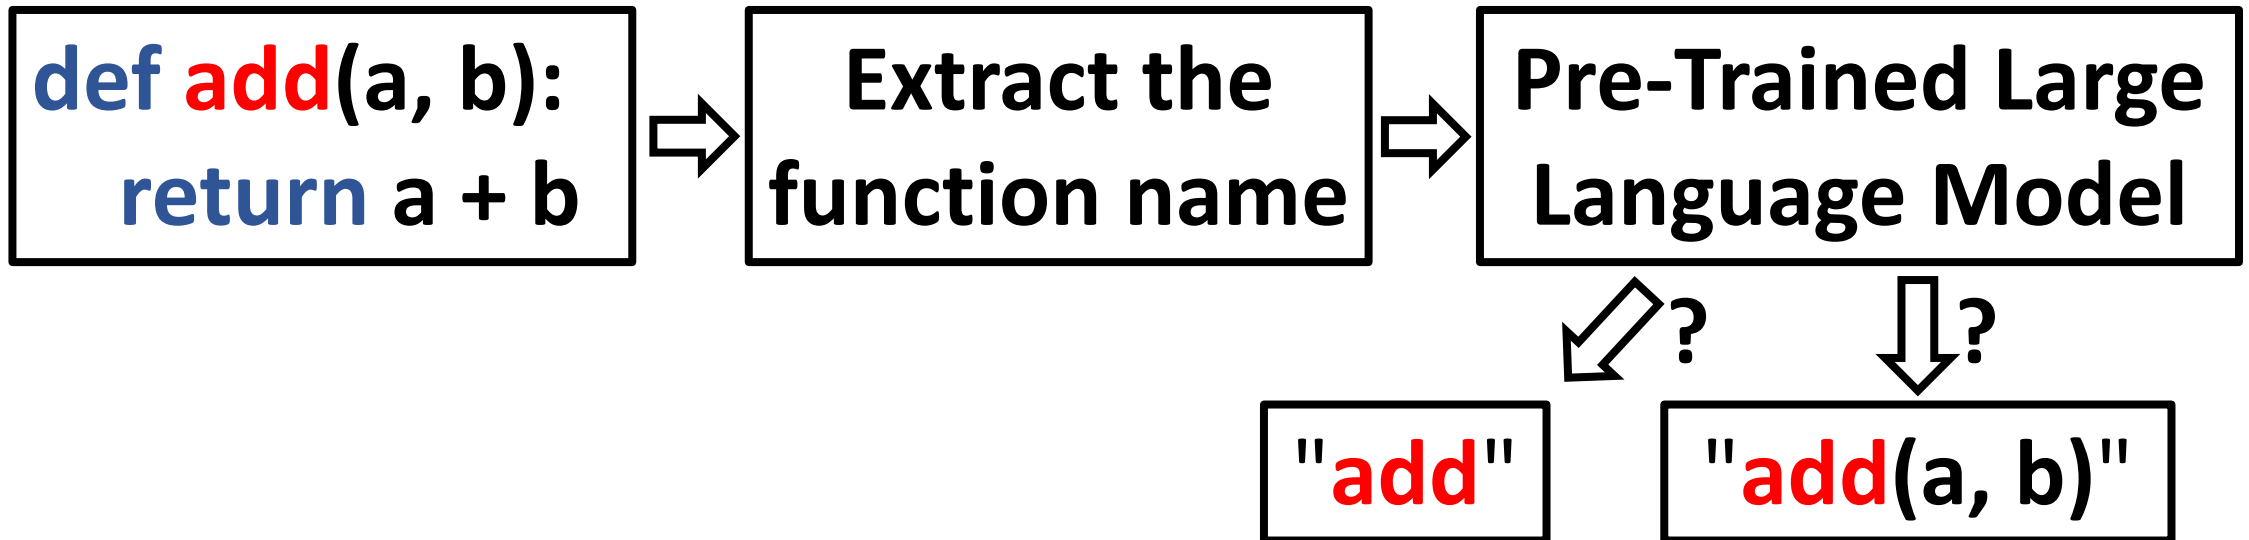
Supervised fine-tuning

- For example, if we want the LLM to extract the function name from the code fragment.



Supervised fine-tuning

- For example, if we want the LLM to extract the function name from the code fragment.



Supervised fine-tuning

- Supervised fine-tuning address this by fine-tuning the pre-trained LLM on a target dataset with **input and output pair**, following the same training process as teacher forcing.

Supervised fine-tuning

- Usually, some **chat templates** and **special tokens** will be added to instruct the LLM.

Supervised fine-tuning

- For example, Llama use <INST> and <SYST> as special tokens and chat templates.

Supervised fine-tuning

- `<<SYS>><</SYS>>` provides context or explanation for the expected role of the LLM.
- `[INST][/INST]` tags are used to encompass entire instructions, such as user questions and inputs.

Supervised fine-tuning

- For example, the input output pair after applying chat template would look like in Llama:

[INST] <<SYS>> You are a helpful assistant in programming. <</SYS>>

Extract the function name from the code.

```
def add(a, b):  
    return a + b [/INST]  
add.
```

Supervised fine-tuning

- For example, the input output pair after applying chat template would look like in Llama:

```
[INST] <<SYS>> You are a helpful assistant in  
programming. <</SYS>>
```

Extract the function name from the code.



Role

```
def add(a, b):  
    return a + b [/INST]  
add.
```

Supervised fine-tuning

- For example, the input output pair after applying chat template would look like in Llama:

[INST] <<SYS>> You are a helpful assistant in programming. <</SYS>>

Extract the function name from the code.

⇒ Input question

```
def add(a, b):  
    return a + b [ /INST]
```

add.

Supervised fine-tuning

- For example, the input output pair after applying chat template would look like in Llama:

[INST] <<SYS>> You are a helpful assistant in programming. <</SYS>>

Extract the function name from the code.

```
def add(a, b):
```

```
    return a + b [/INST]
```

add. \Rightarrow Expected output

Thank you